

ex1

24/25

April 18, 2022

0.1 Exercise sheet 1

Cüneyt Erem 3277992 s6cuerem@uni-bonn.de

Nkeh Victor Ndiwago 3504121 s0vinkeh@uni-bonn.de

Paula Romero Jiménez 3320220 s0parome@uni-bonn.de

ex1)

1. Load the Iris dataset into your notebook from Scikit-Learn

9/9

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

#load iris dataset
from sklearn import datasets
iris = datasets.load_iris()
```

2. Report the descriptive statistics of the features of the iris dataset

- a. Mean, Median, Mode
- b. Variance, MAD, Standard deviation
- c. Quantiles, IQR

```
[ ]: # Dataset preview
print (iris.feature_names)      # Names of features or columns in iris dataset
print (iris.target_names)      # Names of targets in iris dataset
print (iris.data.shape)
```

```
['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
['setosa' 'versicolor' 'virginica']
(150, 4)
```

```
[ ]: # Convert the data into Dataframe
```

```
#iris_df = pd.DataFrame(data= np.c_[iris['data'], iris['target']], columns=
↳ iris['feature_names'] + ['Species'])
iris_df = pd.DataFrame(iris.data, columns=iris.feature_names)
#print(iris_df)
print(iris_df.head())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

```
[ ]: # This function gives us statistical information about the dataset
iris_df.describe()
```

```
[ ]:      sepal length (cm)  sepal width (cm)  petal length (cm)  \
count      150.000000      150.000000      150.000000
mean         5.843333         3.057333         3.758000
std          0.828066         0.435866         1.765298
min          4.300000         2.000000         1.000000
25%          5.100000         2.800000         1.600000
50%          5.800000         3.000000         4.350000
75%          6.400000         3.300000         5.100000
max          7.900000         4.400000         6.900000
```

```
      petal width (cm)
count      150.000000
mean         1.199333
std          0.762238
min          0.100000
25%          0.300000
50%          1.300000
75%          1.800000
max          2.500000
```

a. Mean, Median, Mode

```
[ ]: iris_df.mean()
```

```
[ ]: sepal length (cm)    5.843333
      sepal width (cm)   3.057333
      petal length (cm)  3.758000
      petal width (cm)   1.199333
      dtype: float64
```

```
[ ]: iris_df.median()
```

```
[ ]: sepal length (cm)    5.80
      sepal width (cm)    3.00
      petal length (cm)   4.35
      petal width (cm)    1.30
      dtype: float64
```

```
[ ]: iris_df.mode()
```

```
[ ]:      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
      0                    5.0              3.0              1.4              0.2
      1                    NaN              NaN              1.5              NaN
```

b. Variance, MAD, Standard deviation

```
[ ]: iris_df.var()
```

```
[ ]: sepal length (cm)    0.685694
      sepal width (cm)    0.189979
      petal length (cm)   3.116278
      petal width (cm)    0.581006
      dtype: float64
```

```
[ ]: iris_df.mad()
```

```
[ ]: sepal length (cm)    0.687556
      sepal width (cm)    0.336782
      petal length (cm)   1.562747
      petal width (cm)    0.658133
      dtype: float64
```

We were asking for Median Absolute Deviation.
 sepal length (cm) 1.037822
 sepal width (cm) 0.444781
 petal length (cm) 1.853253
 petal width (cm) 1.037822

```
[ ]: iris_df.std()
```

```
[ ]: sepal length (cm)    0.828066
      sepal width (cm)    0.435866
      petal length (cm)   1.765298
      petal width (cm)    0.762238
      dtype: float64
```

c. Quantiles, IQR

```
[ ]: iris_df.quantile(.75)
```

```
[ ]: sepal length (cm)    6.4
      sepal width (cm)    3.3
      petal length (cm)   5.1
      petal width (cm)    1.8
      Name: 0.75, dtype: float64
```

```
[ ]: iris_df.quantile(.5)
```

```
[ ]: sepal length (cm)    5.80  
      sepal width (cm)    3.00  
      petal length (cm)   4.35  
      petal width (cm)    1.30  
      Name: 0.5, dtype: float64
```

```
[ ]: iris_df.quantile(.25)
```

```
[ ]: sepal length (cm)    5.1  
      sepal width (cm)    2.8  
      petal length (cm)   1.6  
      petal width (cm)    0.3  
      Name: 0.25, dtype: float64
```

```
[ ]: IQR=iris_df.quantile(.75) - iris_df.quantile(.25)  
      print(IQR)
```

```
sepal length (cm)    1.3  
sepal width (cm)     0.5  
petal length (cm)    3.5  
petal width (cm)     1.5  
dtype: float64
```


3. Plot a density plot for each of the variables. Interpret the plots.

```
[ ]: iris_df.rename(columns = {'sepal length (cm)': 'sepal_length', 'sepal width (cm)': 'sepal_width', 'petal length (cm)': 'petal_length', 'petal width (cm)': 'petal_width'}, inplace = True)
```

```
[ ]: list(iris_df)
```


```
[ ]: ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
```

```
[ ]: # Plot for the sepal length  
      iris_df.sepal_length.plot.density(color='green')  
      plt.title('Density plot for sepal length')  
      plt.show()
```



According to the plot, most of the sepals measure around 6 cm long, and if we compare that value with the ones obtained in our mean, median and mode it is correct.


```
[ ]: # Plot for the sepal width  
iris_df.sepal_width.plot.density(color='red')  
plt.title('Density plot for sepal width')  
plt.show()
```




We can observe a clear predominance of 3 cm width, and our values of mean, median and mode are 3, 3.05 and 3.

So we could say in sepals the common points are in one same region, allowing then such a low variance, MAD and standard deviation.

```
[ ]: # Plot for the petal length
iris_df.petal_length.plot.density(color='blue')
plt.title('Density plot for petal length')
plt.show()
```



```
[ ]: # Plot for the petal width
iris_df.petal_width.plot.density(color='orange')
plt.title('Density plot for petal width')
plt.show()
```



In the case of the petals, we have two peaks of density. Therefore, our values of mean, median and mode are not going to coincide with the ones in the plot. Moreover, in petals we observe a higher variance and standard deviation, which is normal taking into account the difference between the common data points.

4. Create a violin plot for the sepal width feature for each class. What can be seen from the plots?

```
[ ]: iris_df2 = pd.DataFrame(data= np.c_[iris['data'], iris['target']], columns=
    ↳ iris['feature_names'] + ['Species'])

iris_df2['Species'].value_counts()

[ ]: 0.0    50
     1.0    50
     2.0    50
     Name: Species, dtype: int64


[ ]: setosa = iris_df.iloc[0:49,1]
     versicolor = iris_df.iloc[50:99,1]
     virginica = iris_df.iloc[100:149,1]
     sepal_w = [setosa, versicolor, virginica]

[ ]: fig, ax = plt.subplots()

     ax.violinplot(sepal_w, positions=None, vert=True, widths=0.5, showmeans=True,
    ↳ showextrema=True, showmedians=False, points=100)

[ ]: {'bodies': [<matplotlib.collections.PolyCollection at 0x7fb10f9cbfd0>,
    <matplotlib.collections.PolyCollection at 0x7fb10f9de2e0>,
    <matplotlib.collections.PolyCollection at 0x7fb10f9de5b0>],
    'cmeans': <matplotlib.collections.LineCollection at 0x7fb1706c3ca0>,
    'cmaxes': <matplotlib.collections.LineCollection at 0x7fb10f9dea00>,
    'cmins': <matplotlib.collections.LineCollection at 0x7fb10f9ded00>,
    'cbars': <matplotlib.collections.LineCollection at 0x7fb113f0a100>}
```


Make sure you use x_label and Y_label!



We observe the difference between the sepal's width of the three species. Setosa is going to have wider sepals than virginica, being versicolor's ones the smallest of the three.

ex2)

8.5/9 points

1) Load the banknote authentication dataset from the given data_banknote_authentication.csv file. How many rows and columns does the dataset contain?

```
[ ]: import pandas as pd
      #import os
      path="data_banknote_authentication.csv"
      dataset=pd.read_csv(path)
      #print(dataset)
      print(dataset.head())
```

2/2

	Variance	Skewness	Curtosis	Entropy	Class
0	NaN	8.6661	-2.8073	-0.44699	0.0
1	4.54590	8.1674	-2.4586	-1.46210	0.0
2	3.86600	-2.6383	1.9242	NaN	0.0
3	3.45660	9.5228	-4.0112	-3.59440	0.0
4	0.32924	-4.4552	4.5718	NaN	0.0

```
[ ]: print(f"there are {len(dataset.index)} rows and {len(dataset.columns)} columns_
      ↪in the dataset")
```

there are 1382 rows and 5 columns in the dataset

1.5/2

2) Mention the different types of variables. Which types does your dataset contain

There are two main types of variables :Categorical variable and Numeric variable Categorical also known as qualitative is subdivided into Nominal(no ordering) and ordinal variable(involves order) while Numerical also known as quantitative variable is subdivided into discrete(takes particular values) and continuous variable(measured on a continuous scale).The dataset contains continuous variables.

Categorical data is qualitative, not quantitative.

3) Count the number of duplicate rows in the dataset. How can you remove the duplicate rows

2/2

```
[ ]: print(dataset.loc[dataset.duplicated(keep="first"),:])
      dataset.duplicated()
      print(f"There are {dataset.duplicated().sum()} duplicated rows")
```

	Variance	Skewness	Curtosis	Entropy	Class
190	0.92970	-3.7971	4.64290	-0.29570	0.0
195	-1.85840	7.8860	-1.66430	-1.83840	0.0
268	0.92970	-3.7971	4.64290	-0.29570	0.0
284	-1.30000	10.2678	-2.95300	-5.86380	0.0
300	0.32920	-4.4552	4.57180	-0.98880	0.0
315	0.32920	-4.4552	4.57180	-0.98880	0.0
345	-1.85840	7.8860	-1.66430	-1.83840	0.0
427	-1.30000	10.2678	-2.95300	-5.86380	0.0
436	0.37980	0.7098	0.75720	-0.44440	0.0
476	0.37980	0.7098	0.75720	-0.44440	0.0
615	-0.20620	9.2207	-3.70440	-6.81030	0.0
691	0.57060	-0.0248	1.24210	-0.56210	0.0
727	-2.64790	10.1374	-1.33100	-5.47070	0.0
1372	NaN	8.6661	-2.80730	-0.44699	0.0
1373	4.54590	8.1674	-2.45860	-1.46210	0.0
1374	3.86600	-2.6383	1.92420	NaN	0.0
1375	3.45660	9.5228	-4.01120	-3.59440	0.0
1376	0.32924	-4.4552	4.57180	NaN	0.0
1377	4.36840	9.6718	-3.96060	-3.16250	0.0
1378	3.59120	3.0129	0.72888	0.56421	0.0
1379	2.09220	-6.8100	NaN	-0.60216	0.0
1380	3.20320	5.7588	-0.75345	-0.61251	0.0
1381	1.53560	9.1772	-2.27180	-0.73535	0.0

There are 23 duplicated rows

```
[ ]: drop= dataset.drop_duplicates()
      #print(drop)
      print(drop.head(), '\n')
      print(drop.info())
```

	Variance	Skewness	Curtosis	Entropy	Class
0	NaN	8.6661	-2.8073	-0.44699	0.0
1	4.54590	8.1674	-2.4586	-1.46210	0.0
2	3.86600	-2.6383	1.9242	NaN	0.0
3	3.45660	9.5228	-4.0112	-3.59440	0.0
4	0.32924	-4.4552	4.5718	NaN	0.0

<class 'pandas.core.frame.DataFrame'>

Int64Index: 1359 entries, 0 to 1371

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	Variance	1217 non-null	float64
1	Skewness	1224 non-null	float64
2	Curtosis	1229 non-null	float64
3	Entropy	1215 non-null	float64
4	Class	1224 non-null	float64

dtypes: float64(5)

memory usage: 63.7 KB

None

4) Count the number of missing values in the dataset

```
[ ]: missing_values= dataset.isna().sum().sum()
      print(f"There are {missing_values} missing values in the dataset.")
```

There are 690 missing values in the dataset. 1/1

5) How an you deal with missing values in a dataset? implement one of the possible methods deleting the rows with a null value for a particular feature replacing the null value with mean meadian or mode assinging a unique category predicting the missing value using algorithms which support missing values

```
[ ]: print('before: ', dataset.isna().sum(), '\n')
      dataset.fillna(dataset.mean().round(1), inplace=True)
      print('after: ', dataset.isna().sum())
```

before: Variance 143
 Skewness 135
 Curtosis 131
 Entropy 146
 Class 135
 dtype: int64

2/2

```
after: Variance    0
Skewness    0
Curtosis    0
Entropy    0
Class    0
dtype: int64
```

ex3) 6.5/7 points Good Job :)

1) Load the dataset from the given dataset.csv file 0.5/1

```
[ ]: import pandas as pd
import seaborn as sns
```

```
[ ]: dataset2 = pd.read_csv("dataset.csv") "Unnamed: 0" column should be removed. It
print(dataset2.head()) is not a feature that should be analysed.
```

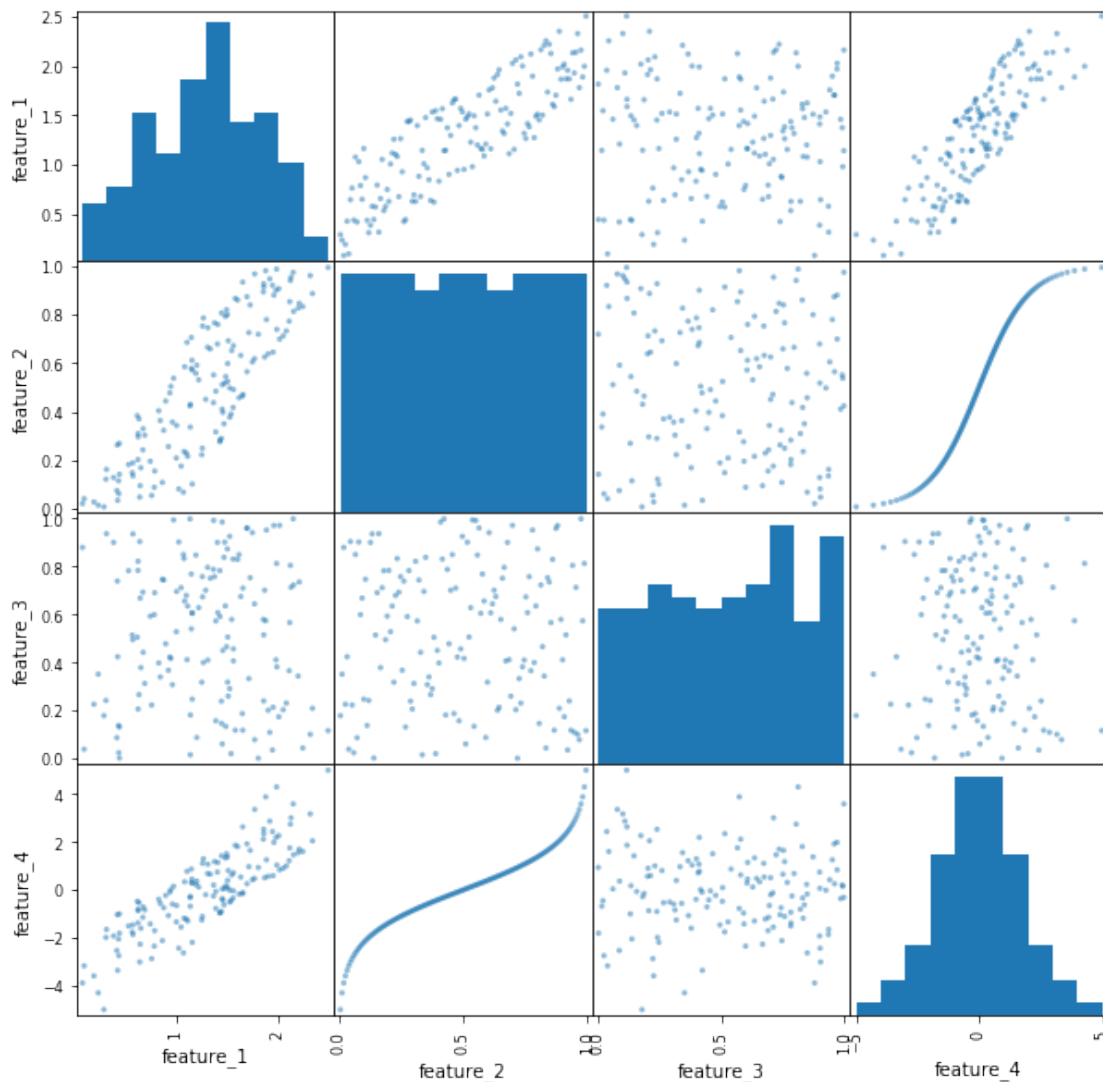
	Unnamed: 0	feature_1	feature_2	feature_3	feature_4
0	0	0.286672	0.006711	0.178739	-4.997212
1	1	0.230586	0.013423	0.351505	-4.297285
2	2	0.074979	0.020134	0.879812	-3.884994
3	3	0.187541	0.026846	0.226149	-3.590439
4	4	0.422490	0.033557	0.424136	-3.360375

2) Plot the scatterplot matrix for the given dataset. What can be seen in the scatterplot matrix? 2/2

```
[ ]: pd.plotting.scatter_matrix(dataset2.iloc[:,1:], figsize=(10,10))
```

```
[ ]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f187e20>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f12d3d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10fa14790>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10fa5dd30>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10fabbe80>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb1706cb5b0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10fae0df0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb111b7d970>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10fa58400>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10fa970a0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f0f1c40>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f0a64f0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f0d0d30>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f0855b0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10f031df0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7fb10efe3670>]])
```

dtype=object)



scatter_matrix shows correlation between different features. feature 1-2 and 1-4 have positive weak correlation, 1-3 has no correlation, 2-4 has positive strong correlation between them.

-
- 3) Which correlation would suit the comparison of feature_1 and feature_3? Calculate the relevant correlation coefficient for the 2 features **2/2**

```
[ ]: dataset_corr = dataset2[['feature_1','feature_3']]

print('pearson: \n', dataset_corr.corr(method='pearson'), '\n')
print('kendall: \n', dataset_corr.corr(method='kendall'), '\n')
print('spearman: \n', dataset_corr.corr(method='spearman'), '\n')
```

```

pearson:
      feature_1  feature_3
feature_1  1.000000 -0.004628
feature_3 -0.004628  1.000000

kendall:
      feature_1  feature_3
feature_1  1.000000 -0.006251
feature_3 -0.006251  1.000000

spearman:
      feature_1  feature_3
feature_1  1.000000 -0.01755
feature_3 -0.01755  1.00000

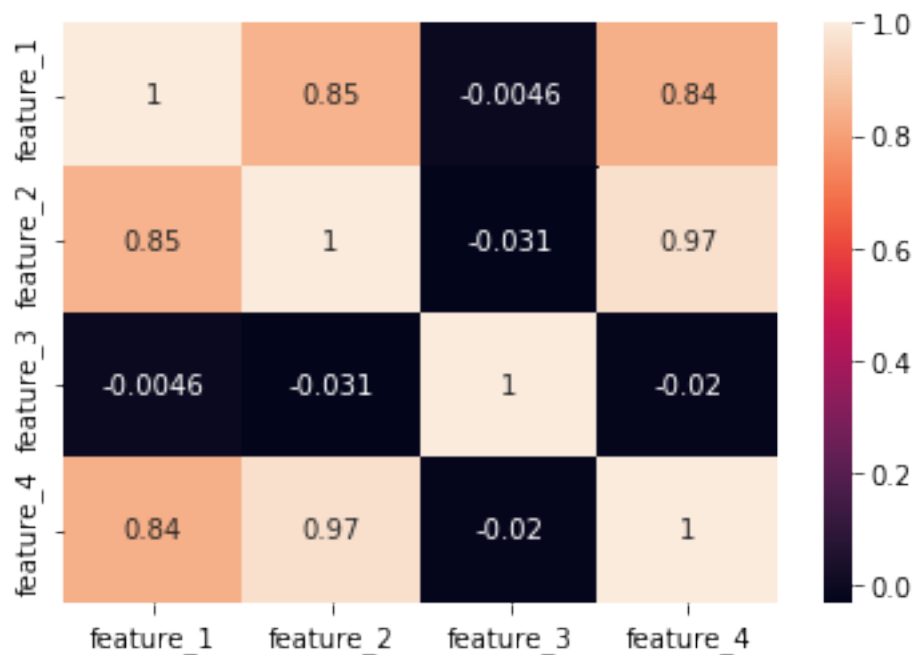
```

there is no obvious correlation between them, and lowest negative correlation is pearson

4) Plot the correlation heatmap of the entire dataset **2/2**

```
[ ]: sns.heatmap(dataset2.iloc[:,1:].corr(method='pearson'), xticklabels=dataset2.
      ↪iloc[:,1:].columns, yticklabels=dataset2.iloc[:,1:].columns, annot=True)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb10efacb50>
```



[]: